

Brian Lester

Toronto, CA

🏠 blester125.com | 📧 blester125 | 🌐 blester125 | 📧 Brian Lester | 📧 Brian Lester

Education

University of Toronto

PH.D. STUDENT WORKING WITH COLIN RAFFEL. ONE YEAR WAS AT THE UNIVERSITY OF NORTH CAROLINA

University of Pittsburgh

DOUBLE MAJOR, 3.51. **COMPUTER SCIENCE** WITH HONORS, 3.86. **NEUROSCIENCE**, 3.19.

Shanghai American School

HIGH SCHOOL, TEAM LEAD IN FIRST ROBOTICS COMPETITION.

Toronto, Canada

2022–Present

Pittsburgh, PA

2012–2016

Shanghai, China

2008–2012

Selected Publications

h-index: 12

The Power of Scale for Parameter-Efficient Prompt Tuning

BRIAN LESTER, RAMI AL-ROUF, NOAH CONSTANT

EMNLP 2021, Online. 5895 citations

Finetuned Language Models Are Zero-Shot Learners

JASON WEI, MAARTEN BOSMA, VINCENT Y ZHAO, KELVIN GUU, ADAMS WEI YU, **BRIAN LESTER**, NAN DU, ANDREW M DAI, QUOC V LE

ICLR 2022, Online. 5302 citations

SPoT: Better Frozen Model Adaptation through Soft Prompt Transfer

TU VU, **BRIAN LESTER**, NOAH CONSTANT, RAMI AL-ROUF, DANIEL MATTHEW CER

ACL 2022, Dublin, Ireland and Online. 341 citations

Scaling Up Models and Data with t5x and seqio

ADAM ROBERTS, HYUNG WON CHUNG, ANSELM LEVSKAYA, GAURAV MISHRA, JAMES BRADBURY, DANIEL ANDOR, SHARAN NARANG, **BRIAN**

LESTER, ET AL. (8/42)

JMLR 2023. 225 citations

Overcoming Catastrophic Forgetting in Zero-Shot Cross-Lingual Generation

TU VU, ADITYA BARUA, **BRIAN LESTER**, DANIEL MATTHEW CER, MOHIT IYER, NOAH CONSTANT

EMNLP 2022, Abu Dhabi, UAE and Online. 77 citations

An Effective Label Noise Model for DNN Text Classification

ISHAN JINDAL, DANIEL PRESSEL, **BRIAN LESTER**, MATTHEW NOKLEBY

NAACL 2019, Minneapolis, Minnesota. 58 citations

The Common Pile v0.1: An 8TB Dataset of Public Domain and Openly Licensed Text

NIKHIL KANDPAL*, **BRIAN LESTER***, COLIN RAFFEL*, ET AL. (2/27)

NeurIPS 2025; Datasets & Benchmarks, San Diego, CA. 24 citations

Git-Theta: A Git Extension for Collaborative Development of Machine Learning Models

NIKHIL KANDPAL*, **BRIAN LESTER***, MOHAMMED MUQEETH, ANISHA MASCARENHAS, MONTY EVANS, VISHAL BASKARAN, TENGHAO HUANG,

HAOKUN LIU, COLIN RAFFEL

ICML 2023, Honolulu Hawaii. 18 citations

A Library for Rapid Modeling, Experimentation and Development of Deep Learning Algorithms targeting NLP

DANIEL PRESSEL, SAGNIK RAY CHOUDHURY, **BRIAN LESTER**, YANJIE ZHAO, MATT BARTA

ACL 2018; NLP-OSS Workshop, Melbourne, Australia. 16 citations

Realistic Evaluation of Model Merging for Compositional Generalization

DEREK TAM*, YASH KANT*, **BRIAN LESTER***, IGOR GILITSCHENSKI, COLIN RAFFEL

TMLR 2026, 2026. 15 citations

Training LLMs over Neurally Compressed Text

BRIAN LESTER, JAEHOON LEE, ALEX ALEMI, JEFFREY PENNINGTON, ADAM ROBERTS, JASCHA SOHL-DICKSTEIN, NOAH CONSTANT

TMLR 2024. 17 citations

Reducing Retraining by Recycling Parameter-Efficient Prompts

BRIAN LESTER*, JOSHUA YURTSEVER*, SIAMAK SHAKERI, NOAH CONSTANT

Preprint arXiv:2208.05577, 2022. 14 citations

Multiple Word Embeddings for Increased Diversity of Representation

BRIAN LESTER, DANIEL PRESSEL, AMY HEMMETER, SAGNIK RAY CHOUDHURY, SRINIVAS BANGALORE

Preprint arXiv:2009.14394, 2020. 11 citations

Constrained Decoding for Computationally Efficient Named Entity Recognition Taggers

BRIAN LESTER, DANIEL PRESSEL, AMY HEMMETER, SAGNIK RAY CHOUDHURY, SRINIVAS BANGALORE

EMNLP 2020; Findings, Online. 10 citations

iobes: Library for Span Level Processing

BRIAN LESTER

ACL 2020; NLP-OSS Workshop, Online. 8 citations

TokSuite: Measuring the Impact of Tokenizer Choice on Language Model Behavior

GÜL SENA ALTINTAŞ, MALIKEH EHGAGHI, **BRIAN LESTER**, FENGYUAN LIU, WANRU ZHAO, MARCO CICCONE, COLIN RAFFEL

ICML 2026, 2026. 3 citations

Intent Features for Rich Natural Language Understanding

BRIAN LESTER, SAGNIK RAY CHOUDHURY, RASHMI PRASAD, SRINIVAS BANGALORE

NAACL 2021; Industry Track, Online.

Work Experience

Google DeepMind

SENIOR RESEARCH ENGINEER

Mountain View, California

2020–Present

Deep Learning research with a focus on Natural Language Processing, large pre-trained models, and zero-shot transfer.

- Prompt Tuning: an efficient method of controlling large frozen pre-trained language models based on T5. Matches performance of full fine-tuning using only 0.003% of the parameters. Open-sourced our codebase and it has enabled 5 published papers, 1 product launch, and at least 3 more in-flight papers.
- Flan: Multitask training for a 137 billion parameter transformer-based decoder-only language model to create a model that is more effective at zero-shot prompting and performs better using Prompt Tuning.
- SPoT: Using multitask prompts as strong initialization for Prompt Tuning resulting in increased performance. Also used prompt similarity to estimate task similarity and to predict transferability.
- Added partial network training, lazy loading, and pre-filling of the auto-regressive cache to t5x, the open-source reimplementation of T5 in Jax. This final change reduced inference latency from 30 seconds to 2.4.

Interactions

Ann Arbor, Michigan

MACHINE LEARNING ENGINEER

2018–2020

Built production grade deep learning solutions and lead research efforts to push the boundaries of performance.

- Designed novel neural network architectures for calibrated intent detection, slot filling, and named entity linking using ConvNets, bLSTM-CRFs, ranking models, and transformer-based seq2seq models.
- Designed label space, annotation guidelines, and data collection method for NLU component of dialogue systems.
- Created a cloud-native model training platform based on declarative pipelines and kubernetes. Built a deployment platform that powers NLU for multiple production dialogue systems.
- Built efficient, batched implementations of complex neural network architectures such as Beam Search. My CRF implementation reduced training time by a factor of 10.

Trove

Ann Arbor, Michigan

LEAD MACHINE LEARNING RESEARCH ENGINEER

2017–2018

Created a model training and serving platform that processed 200 million emails per day. Provided technical leadership to the ML team.

- Designed ConvNets for text classification to find sentences that contain questions. This powered a user-facing feature and was used to featurize the social graph created from email.
- Created neural ranking model was used to find coreferent mentions in the text and provide context to users.
- Used lexical features, as well as connectivity information in the email social graph, to identified bot accounts.

Visteon Corporation

Van Buren, Michigan

SOFTWARE ENGINEERING INTERN

2015

Designed an adaptive system to minimize voice recognition errors based on ASR confidence scores. We patented this system and it is used in Mazda Cars.

Presentations

Prompt Tuning

University of Michigan

University of North Carolina

2021–2022

DEEP LEARNING

An overview of my work on Prompt Tuning, as well as our work—Flan and SPoT—directly built on Prompt Tuning. The talk includes a collection of insights about the behavior of soft prompts aggregated from others' followup work.

NeurIPS Spotlight Talk on Mead-Baseline

NeurIPS OSS Workshop

DEEP LEARNING

December 2018

A spotlight talk, at the Open Source Software workshop at NeurIPS 2018, about our open-source toolkit, Mead-Baseline.

Confidence and Calibration of Neural Network Models

EMU ML Conference

DEEP LEARNING

March 2020¹

An overview of techniques used to adjust model calibration, evaluation of models that have the ability to “reject” decision with low confidence, and their uses in the NLU unit of a production dialogue system.

Padding in Neural Networks for Natural Language Processing

A2D-NLP

NATURAL LANGUAGE PROCESSING

February 2020

A survey of NLP building blocks with a focus on correctness and the need for padding in complex situations as well as places it is unexpected, like max-pooling following a 1D convolution.

Optimization via NumPy and Cython

Michigan Python Meetup

NUMERICAL COMPUTATION

January 2020

I use a series of optimizations for computing pairwise Manhattan distance to introduce core NumPy concepts and Cython to reduce the runtime from multiple hours to just seconds.

Input Representations of Deep Neural Networks

PyData Ann Arbor

DEEP LEARNING

October 2017

Using learned character-compositional input representations to create Deep Neural Networks with an open vocabulary.

¹Canceled due to COVID-19

Skills

Deep Learning

Extensive experience building novel Neural Network architectures, generally for NLP. High-performance training with Data and Model Parallelism, including multihost distributed training on TPU.

Infrastructure

Build and deploy with Kubernetes, Docker, Flux, MongoDB, Apache Nifi, Github Actions, and GitLab CI/CD. Experience training large neural networks on Google Cloud (GCP)

Toolkits

Jax, Flax, PyTorch, Tensorflow, NumPy, Pandas, Faiss, SpaCy, NLTK, Tensorflow-Datasets, Seaborn, and Matplotlib.

Languages

Python, Cython, Java, C, Javascript, C++, Elisp, and $\text{\LaTeX}$.

Service & Public Scholarship

REVIEWING

- 2023–2025 NeurIPS
- 2022 IEEE Transactions on Affective Computing
- 2022 ARR: ACL Rolling Review
- 2022–2024 NAACL
- 2025 ACL
- 2025 EMNLP
- 2025 ICLR
- 2025 ICML
- 2020–2021 Computer Speech and Language
- 2019 CoNLL

PUBLIC SCHOLARSHIP

- 2020¹ **a²-dlearn**: Helped organize logistics, recruit speakers, and acquire funding through sponsorships