

Brian Lester

2251 Glenwood St. Trenton, MI. 48183

☎ (734) 834-2546 | ✉ blester125@gmail.com | 🌐 blester125.com | 📱 blester125 | 📺 blester125 | 🎓 Brian Lester

Work Experience

Interactions

Ann Arbor, Michigan

MACHINE LEARNING ENGINEER

2018–Present

Driving Deep Learning adoption. Built, maintained, designed, and productionized deep learning models to facilitate customer care interactions. Used for applications like dialogue systems and the retrieval, ranking, and categorization of social media posts (Twitter, Facebook, and web forums) to increase the effectiveness of customer service representatives.

- **Mead-Baseline:** I improve and maintain Mead-Baseline—our open-source, deep-learning toolkit and the sole method for training and deploying deep learning models in our entire company. Mead-Baseline doubles as our research platform and because of its extensible design I have used it for both novel research and to reproduce and vet current literature. Mead-Baseline includes my performant implementation of a CRF. My CRF is ten times faster than our original implementation making it undoubtedly one of the fastest ways to train a bLSTM-CRF for sequence tagging. My CRF also yields strong performance by helping a user inject constraints into the model as dictated by their tagging scheme. I have also accelerated the implementations of other complex modules such as Beam Search.
- **NLU Annotation Design:** Designed the label space for the Natural Language Understanding module of a customer self service dialogue system in the technical support domain. My design focused on using general intents, complex entities, and relations to cover a diverse and complex conversation space.
- **Deep Learning Training Platform:** Designed and created our cloud-native, model-training platform used by the whole company to train deep models. Training pipelines are built declaratively by specifying a directed acyclic graph where each node represents some unit of work to be run inside a Docker container. This graph is then executed on our Kubernetes cluster. Using Kubernetes helps us effectively manage the usage of shared GPUs as well as schedule many jobs in parallel allowing Hyper Parameter Optimization at scale. This platform enables building complex multi-step pipelines (often one step is training a model via Mead-Baseline) that transform raw data into a model that is ready for production.
- **Deep Learning Deployment Platform:** Designed and implemented our model server. This server, backed by TensorFlow Serving, enables rich Natural Language Understanding via cascading calls to a series of deep learning models. Both the model server itself and the TensorFlow Serving backend are deployed via Kubernetes. This model server is currently powering NLU for several production dialogue systems.
- **Production Models:** Created a wide range of machine learning models that are used in production. Model architectures and tasks range from ConvNets for classification and intent detection, bLSTM-CRF taggers for general NER and client-specific slot filling, mention-ranking models for relation extraction, to Transformer-based seq2seq models that generate suggested agent responses. Production models were also calibrated to produce reliable scores via post-hoc methods and secondary confidence models.

Trove

Ann Arbor, Michigan

LEAD MACHINE LEARNING RESEARCH ENGINEER

2017–2018

Replaced heuristics with Machine Learning to enhance features. Created a model training and serving platform that processed 200 million emails a day. Provided technical leadership to the data science team.

- **Question Detection:** Text classification with a ConvNet to find sentences that contain questions. The presence of questions in an email was used as a feature in our email social graph.
- **Coreference Resolution:** A neural ranking model was used to find coreferent mentions in the text and provide context to users.
- **Bot detection:** Using lexical features, as well as connectivity information in the email social graph, we identified bot accounts.

Visteon Corporation

Van Buren, Michigan

SOFTWARE ENGINEERING INTERN

2015

Optimized Voice Recognition in Mazda Cars. Designed an adaptive system to minimize voice recognition errors. We patented this system (US 9,984,688 B2) and it is used in production for Mazda Cars.

Skills

Machine and Deep Learning

Extensive experience using Pytorch, Tensorflow, and Scikit-learn for NLP applications.

Infrastructure

Building and deploying with Kubernetes, Docker, Flux, Prometheus, MongoDB, and GitLab CI/CD.

Toolkits

NumPy, Pandas, Faiss, Gensim, SpaCy, NLTK, SciPy, and Matplotlib.

Languages

Python, Cython, Java, C, Javascript, C++, and \LaTeX .

Projects

Text Rank

Personal Project

NATURAL LANGUAGE PROCESSING

An implementation of Text Rank in Python that reproduces the results from the original paper. I currently use it as a vehicle for personal research: investigating if deep learning based sentence similarities will yield better summaries. Results forthcoming.

String Distance

Open Source Library

NATURAL LANGUAGE PROCESSING

A collection of various minimum edit distance algorithms as well as token based methods like Jaccard overlap. The algorithms are implemented in Cython and scale well. The library is able to compute minimum edit distances between entire Wikipedia pages in under a second.

Quick KNN

Open Source Library

INFORMATION RETRIEVAL

Implementations of Locality Sensitive Hashing. Supports using MinHash to approximate Jaccard similarity and Random Hyperplanes for a cosine based LSH. This library enables users to find similar items in very large corpora quickly.

Dependency Parsing

Personal Project

NATURAL LANGUAGE PROCESSING

Dependency parsing via Deep Learning models written in PyTorch. Supports training parsers via dynamic oracles using either the Arc Eager or the Arc Hybrid transition scheme or a Graph based parser using Biaffine Attention. Currently working on creating a shared interface between the two parser types with plans to release it as an open source library.

Decomposable Attention

Personal Project

NATURAL LANGUAGE PROCESSING

A reimplementation of the paper "A Decomposable Attention Model for Natural Language Inference" in DyNet. I was able to reproduce the paper results via extensive hyper parameter tuning.

Steering Angles

Personal Project

DEEP LEARNING

Used a deep convolutional neural network written in TensorFlow to create a steering agent for an autonomous car. Reduced training time of network using architecture optimizations like residual connections and batch normalization.

Multi-Digit Recognition

Personal Project

DEEP LEARNING

Multi-Digit sequence recognition in natural scenes via a deep, convolutional neural network written with TensorFlow. The end-to-end network allowed optimization of the whole process rather than dividing it into localization, segmentation, and recognition tasks.

RFC 793 Transmission Control Protocol

University of Pittsburgh

COMPUTER NETWORKING

Fall 2015

A full implementation of RFC 793 TCP written in C++. It insures reliable communication between two hosts and supports multiple data packets in flight using a "Go Back N" strategy.

Particle Simulation

University of Pittsburgh

HIGH PERFORMANCE COMPUTING

Spring 2015

Simulation of particle interactions written in MPI. The simulation was parallelized via a ring algorithm and was run on the Stampede super computer.

Secure File Sharing System

University of Pittsburgh

APPLIED CRYPTOGRAPHY AND NETWORK SECURITY

Spring 2016

A distributed, group-based file sharing system that is secured using cryptographic techniques such as symmetric and public key cryptography, ephemeral key exchange, digital signatures, and two factor authentication.

Presentations

NeurIPS Spotlight Talk on Mead-Baseline

NeurIPS OSS Workshop

DEEP LEARNING

December 2018

A spotlight talk, at the Open Source Software workshop at NeurIPS 2018, highlighting the advantages of using our open-source, model-building toolkit, Mead-Baseline, which provides high-level model abstractions, correct evaluation metrics, and fast runtimes. Mead-Baseline is the foundation of all Deep Learning work, in both research and production, at Interactions.

Confidence and Calibration of Neural Network Models

EMU ML Conference

DEEP LEARNING

March 2020,¹

In most settings a model is trained and evaluated on a held out test set; however, in some domains one would rather have no answer than an uncertain one. Modern Neural Networks are notoriously miscalibrated, that is, they are far too confident in their predictions. This talk discusses techniques used to evaluate model calibration and rejection of examples based on confidence as well as a summary of current work to produce high-fidelity confidence scores in the NLU component of a real world Dialogue System.

Padding in Neural Networks for Natural Language Processing

A2D-NLP

NATURAL LANGUAGE PROCESSING

February 2020

A deep dive into the implementation of many NLP building blocks focusing on the subtleties of the padding and masking needed to correctly process batches. It addressed the masking needed for operations like token level losses, mean pooling, and attention; complex operations such as the CRF forward algorithm and Viterbi Decoding; and subtle operations that you might not expect to need masking like max pooling following a 1D convolution.

Optimization via NumPy and Cython

Michigan Python Meetup

NUMERICAL COMPUTATION

January 2020

Framed around the problem of calculating the pairwise Manhattan distance between a collection of points, this talk walks through a series of optimizations that introduce core NumPy concepts like Broadcasting and Vectorization before using a bespoke Cython implementation to scale up to 10s of thousands of points in just seconds compared to the multi hour runtime of a pure Python solution.

a2-dlearn 5

Ann Arbor, Michigan

DEEP LEARNING

2020,¹

Helped organize, recruit speakers, and acquire funding through sponsorships for the premier deep learning event in Michigan. Attendance regularly breaks 100 people and this year is expected to be bigger than ever.

Input Representations of Deep Neural Networks

PyData Ann Arbor

DEEP LEARNING

October 2017

A talk on creating robust models via learned character compositional input representations based on bLSTMs to support an open vocabulary for Deep Neural Networks.

Conference Publications

An Effective Label Noise Model for DNN Text Classification

I. JINDAL, D. PRESSEL, B. LESTER, M. NOKLEBY

Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers), 2019, Minneapolis, Minnesota

Workshop Publications

Baseline: Strong, Extensible, Reproducible, Deep Learning Baselines for NLP

D. PRESSEL, B. LESTER, S. R. CHOUDHURY, M. BARTA, Y. ZHAO, A. HEMMETER

Neural Information Processing Systems Open Source Software Workshop (2018). 2018

Baseline: A Library for Rapid Modeling, Experimentation and Development of Deep Learning Algorithms targeting NLP

D. PRESSEL, S. RAY CHOUDHURY, B. LESTER, Y. ZHAO, M. BARTA

Proceedings of Workshop for NLP Open Source Software (NLP-OSS), 2018, Melbourne, Australia

¹ Canceled due to COVID-19

Other Publications

Computationally Efficient NER Taggers with Combined Embeddings and Constrained Decoding

B. LESTER, D. PRESSEL, A. HEMMETER, S. R. CHOUDHURY

arXiv preprint arXiv:2001.01167 (2020). 2020

Dynamically Adjusting a Voice Recognition System

B. LESTER, S. M. PANAINTE

US Patent 9,984,688, 2018

Education

University of Pittsburgh

Pittsburgh, PA

DOUBLE MAJOR, 3.51. **COMPUTER SCIENCE** WITH HONORS, 3.86. **NEUROSCIENCE**, 3.19.

2012–2016

Shanghai American School

Shanghai, China

HIGH SCHOOL

2008–2012

Advanced Placement Curriculum. Team lead for FIRST Robotics Competition.